



ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA

# APAI Lab01: PULP Embedded Programming

**Davide Nadalini, Lorenzo Lamberti, Alberto Dequino,  
Luca Bompani, Francesco Conti.**

(University of Bologna)

`d.nadalini@unibo.it`

`alberto.dequino@unibo.it`

# Objective of the Class

**Intro:** PULP platform and the PULP-SDK

**Tasks:** some basics of C programming on PULP:

- Hello world
- vector sum,
- Matrix-vector multiplication
- Measuring execution performance

**Programming Language:** C

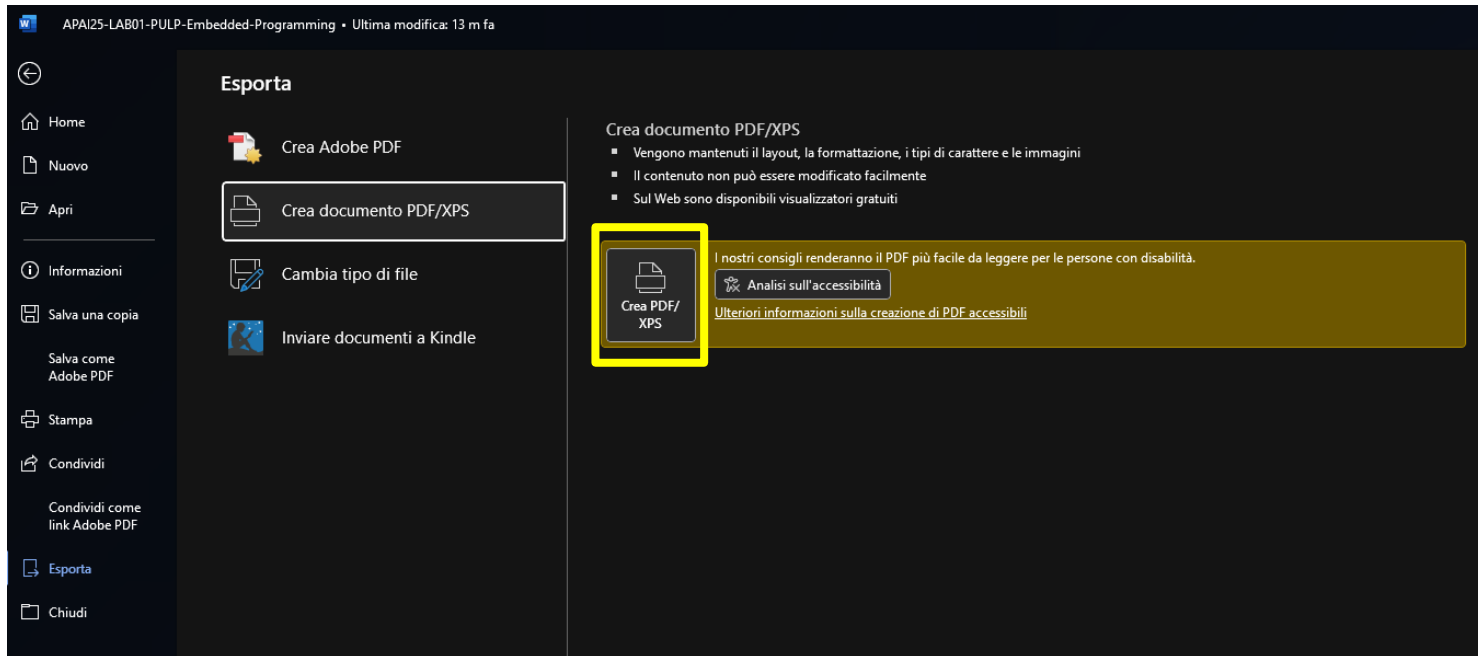
**Lab duration:** 3h

The class is meant to be interactive: coding together and on your own!

# How to deliver the Assignment

You will deliver ONLY the PDF assignment, no code

- Download the assignment file from Virtuale.
- Fill the results required by the assignment.
- Export to pdf format.
- Rename the file to: LAB<number\_of\_the\_lesson>\_APAI\_<your\_name>.pdf
- Use Virtuale platform to load ONLY your .pdf file

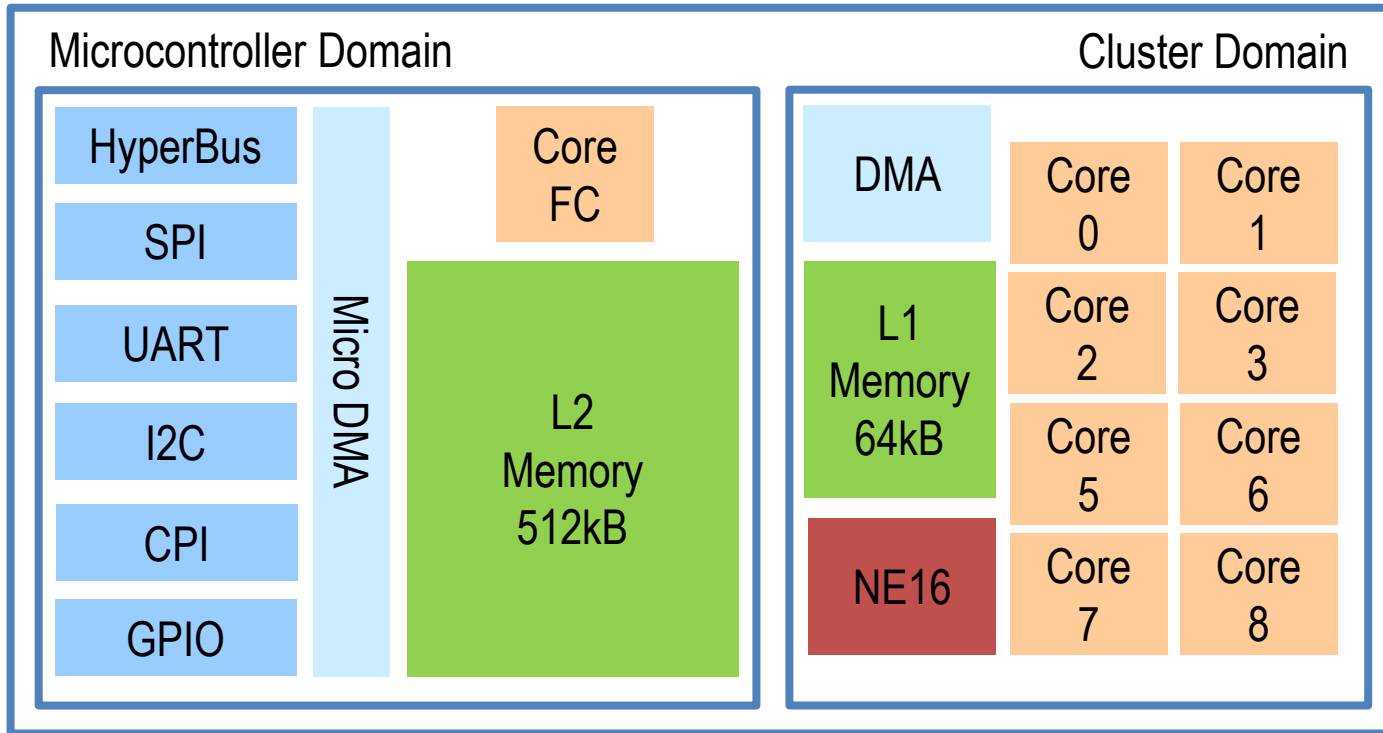




ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA

# INTRO

# Get confident with the PULP Platform



- **Cores:** 1 + 8
- **On-chip Memories**
  - A level 2 Memory, shared among all cores
  - A level 1 Memory, shared by the 8-cores cluster
- **cluster-DMA:** A multi-channel 1D/2D DMA, controlling the transactions between the L2 and L1 memories
- **micro-DMA:** A smart, lightweight and completely autonomous DMA () capable of handling complex I/O scheme
- **Bus+Peripherals:** HyperBus, I2S, CPI, timers, SPI, GPIOs, etc...

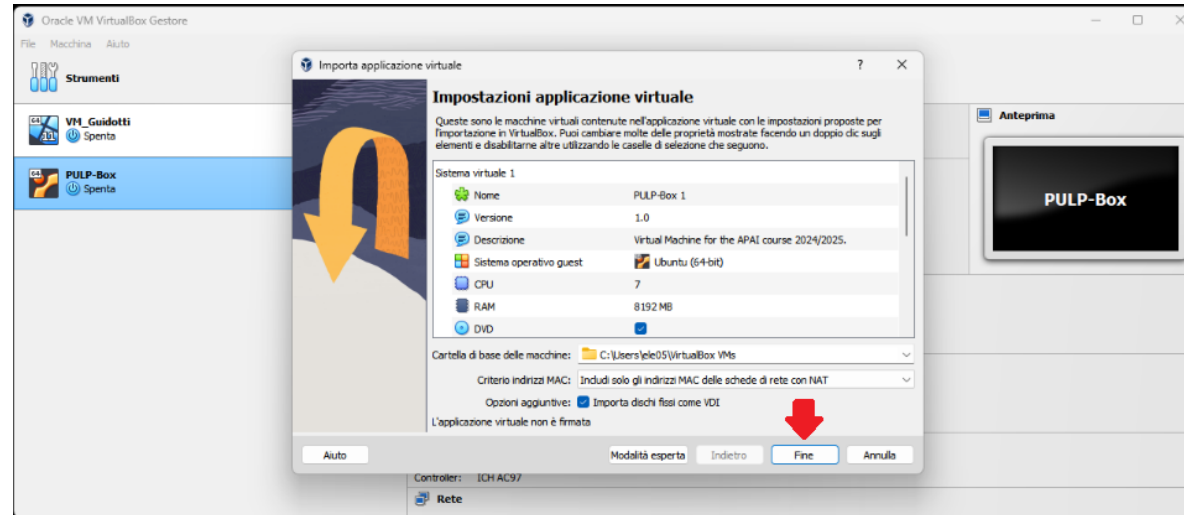
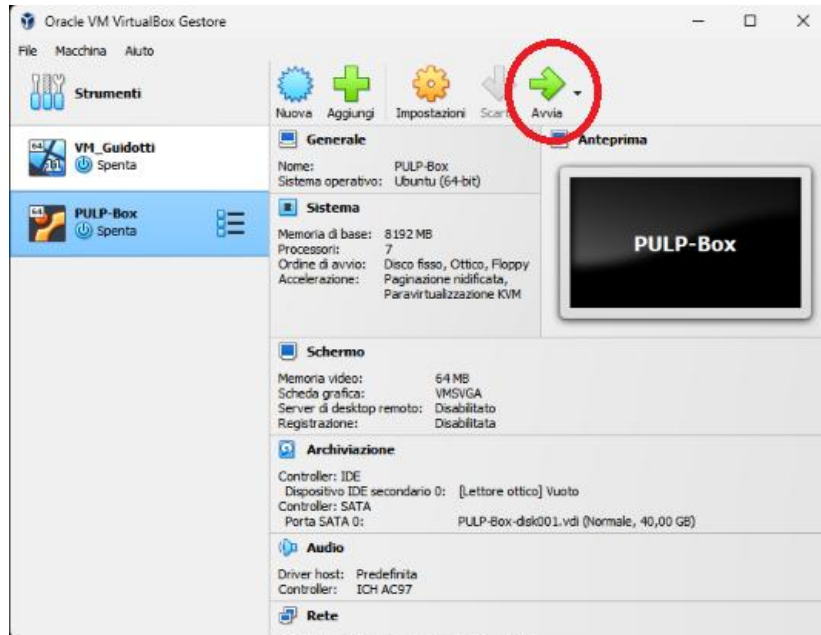
**GitHub HW Project:** <https://github.com/pulp-platform/pulp>  
**HW Documentation:** <https://raw.githubusercontent.com/pulp-platform/pulp/master/doc/datasheet.pdf>

**NB: this is the architecture you find on our nano-drones and GAP boards!**



# Opening the VM and VSCode

1. On the lab's PCs, open the file explorer and go to This PC, C:/VM\_Nadalini/
2. Double click on PULP-box.ova
3. VirtualBox opens, just click on “Fine”
4. Wait for the VM to be imported
5. Open the VM with “Avvia”



Password is 'pulp'

# Opening the Docker with VSCode


1. Open a terminal (right click – open a new terminal)
2. From the terminal, open VSCode in the folder of the Docker

```
$ cd APAI-Docker  
$ code .
```

1. Reopen the APAI-Docker folder in VSCode (click on “Reopen in container”)
1. Now you can use the **integrated terminal (open with CTRL+J)** to run your applications!

**IMPORTANT:** every time you open a **new terminal** to work on PULP, launch

```
$ source setup-pulp-sdk.sh
```



# Getting Started: *Helloworld*

**IMPORTANT:** activate the pulp-sdk module file every time a new shell is open.

```
$ source setup-pulp-sdk.sh
```

## HOW TO RUN THE CODE:

```
$ git clone https://github.com/EEESlab/APAI25-LAB01-PULP-Embedded-Programming
$ cd APAI25-LAB01-PULP-Embedded-Programming
$ cd pulp-helloworld/
$ make clean all run
```

*test.c*

- Can you see the **Helloworld** from **PULP!** ?

```
int main()
{
    printf("Helloworld from PULP!\n");
}
```



# Behind the box: Build Automation with Makefiles

**Clean** previous build processes  
( folder *BUILD/* )

**Compile** the program  
and produce the **binary code**

Run **binary code** on the  
target platform

```
$ make clean all run
```

**Build automation** is the process of automating the main steps required to create a software, including *compiling*, *assembling*, *linking* and (possibly) *testing*

**Make** is one of the most widespread utilities

➤ configuration files are called **Makefiles**

A *Makefile* contains rules in the form:

target: prerequisites

<TAB> command

*Makefile*

```
APP = test

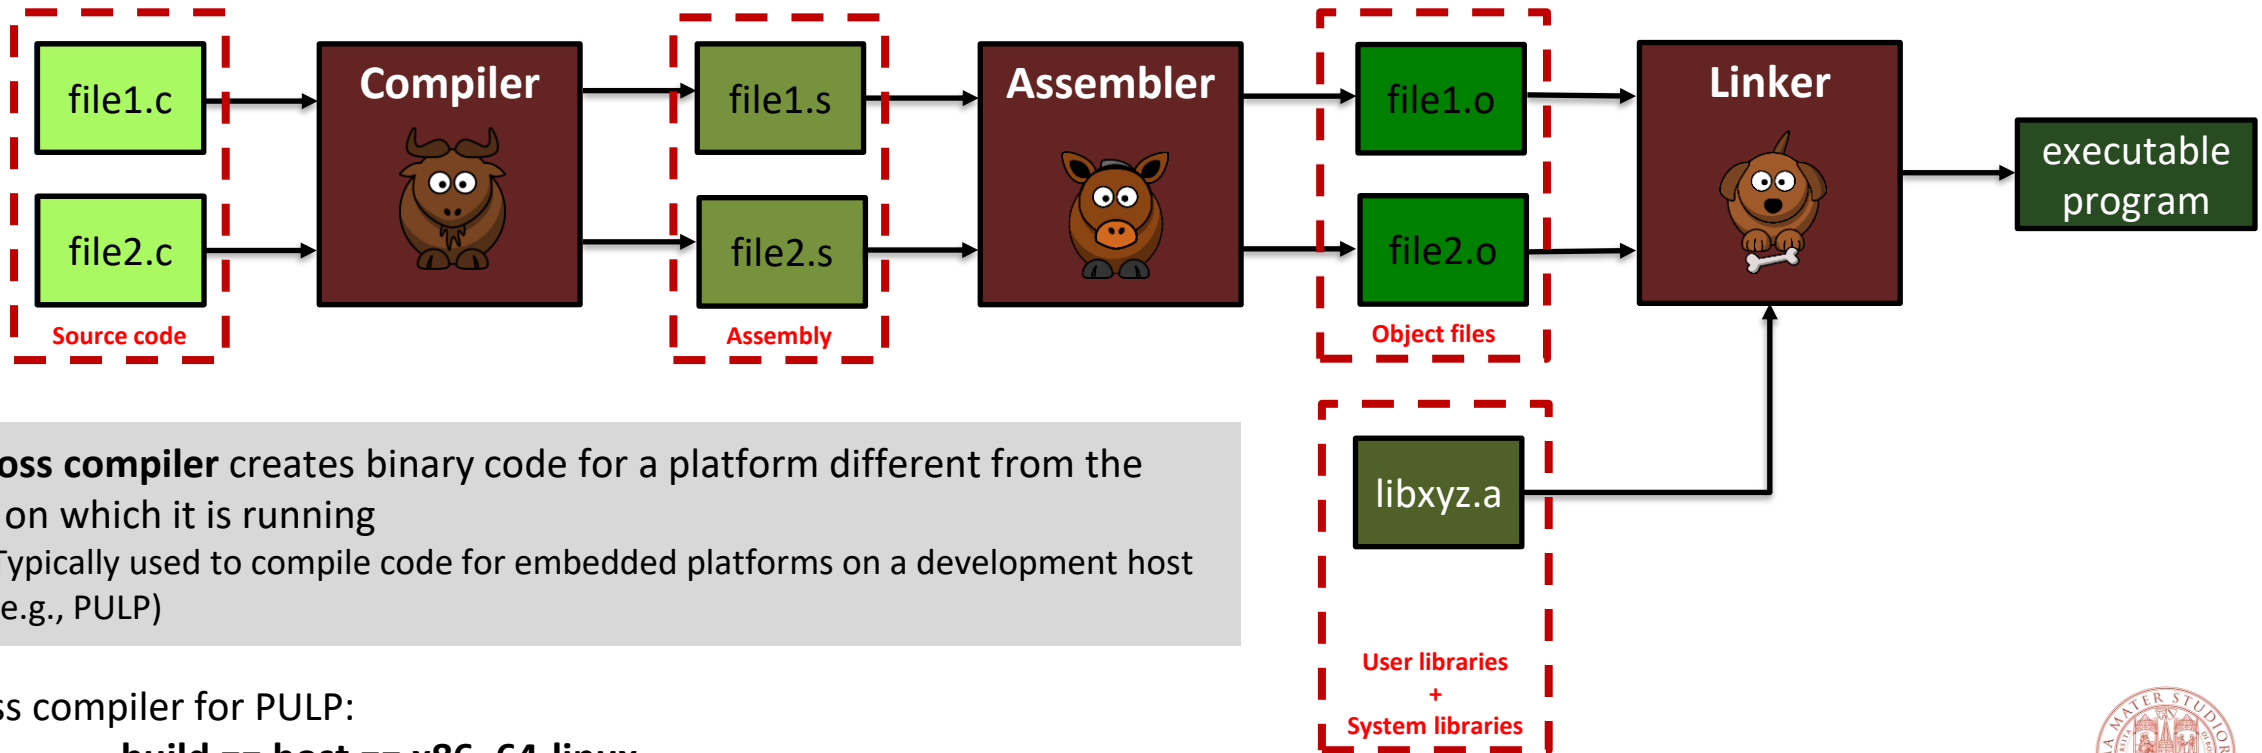
# This is a comment
APP_SRCS += test.c
APP_CFLAGS += -O3 -g
APP_LDFLAGS +=

include $(RULES_DIR)/pmsis_rules.mk
```

# Compilation toolchain

A compilation toolchain includes several tools to achieve its final goal

- The **COMPILER** translates high-level source code (e.g., C) to a lower-level representation (e.g., assembly)
- The **ASSEMBLER** is program that translates assembly language to machine language
- The **LINKER** combines multiple object files into a single executable



A **cross compiler** creates binary code for a platform different from the one on which it is running

- Typically used to compile code for embedded platforms on a development host (e.g., PULP)

Cross compiler for PULP:

**build == host == x86\_64-linux**

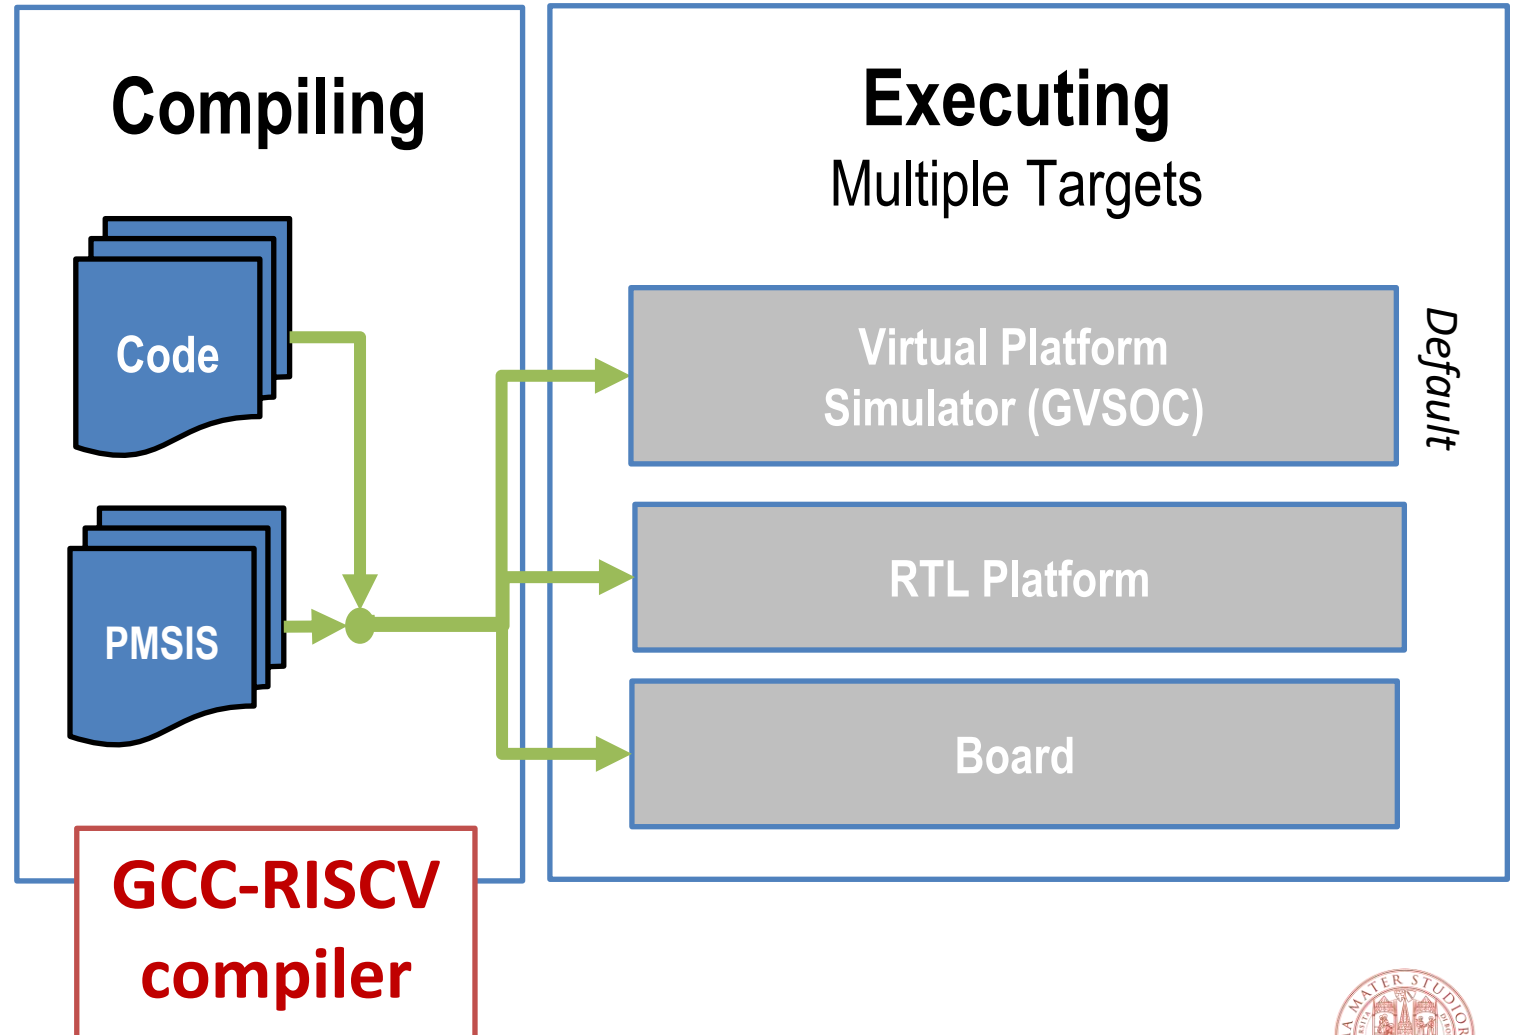
**target = riscv-none**

# PULP Software Environment Workflow

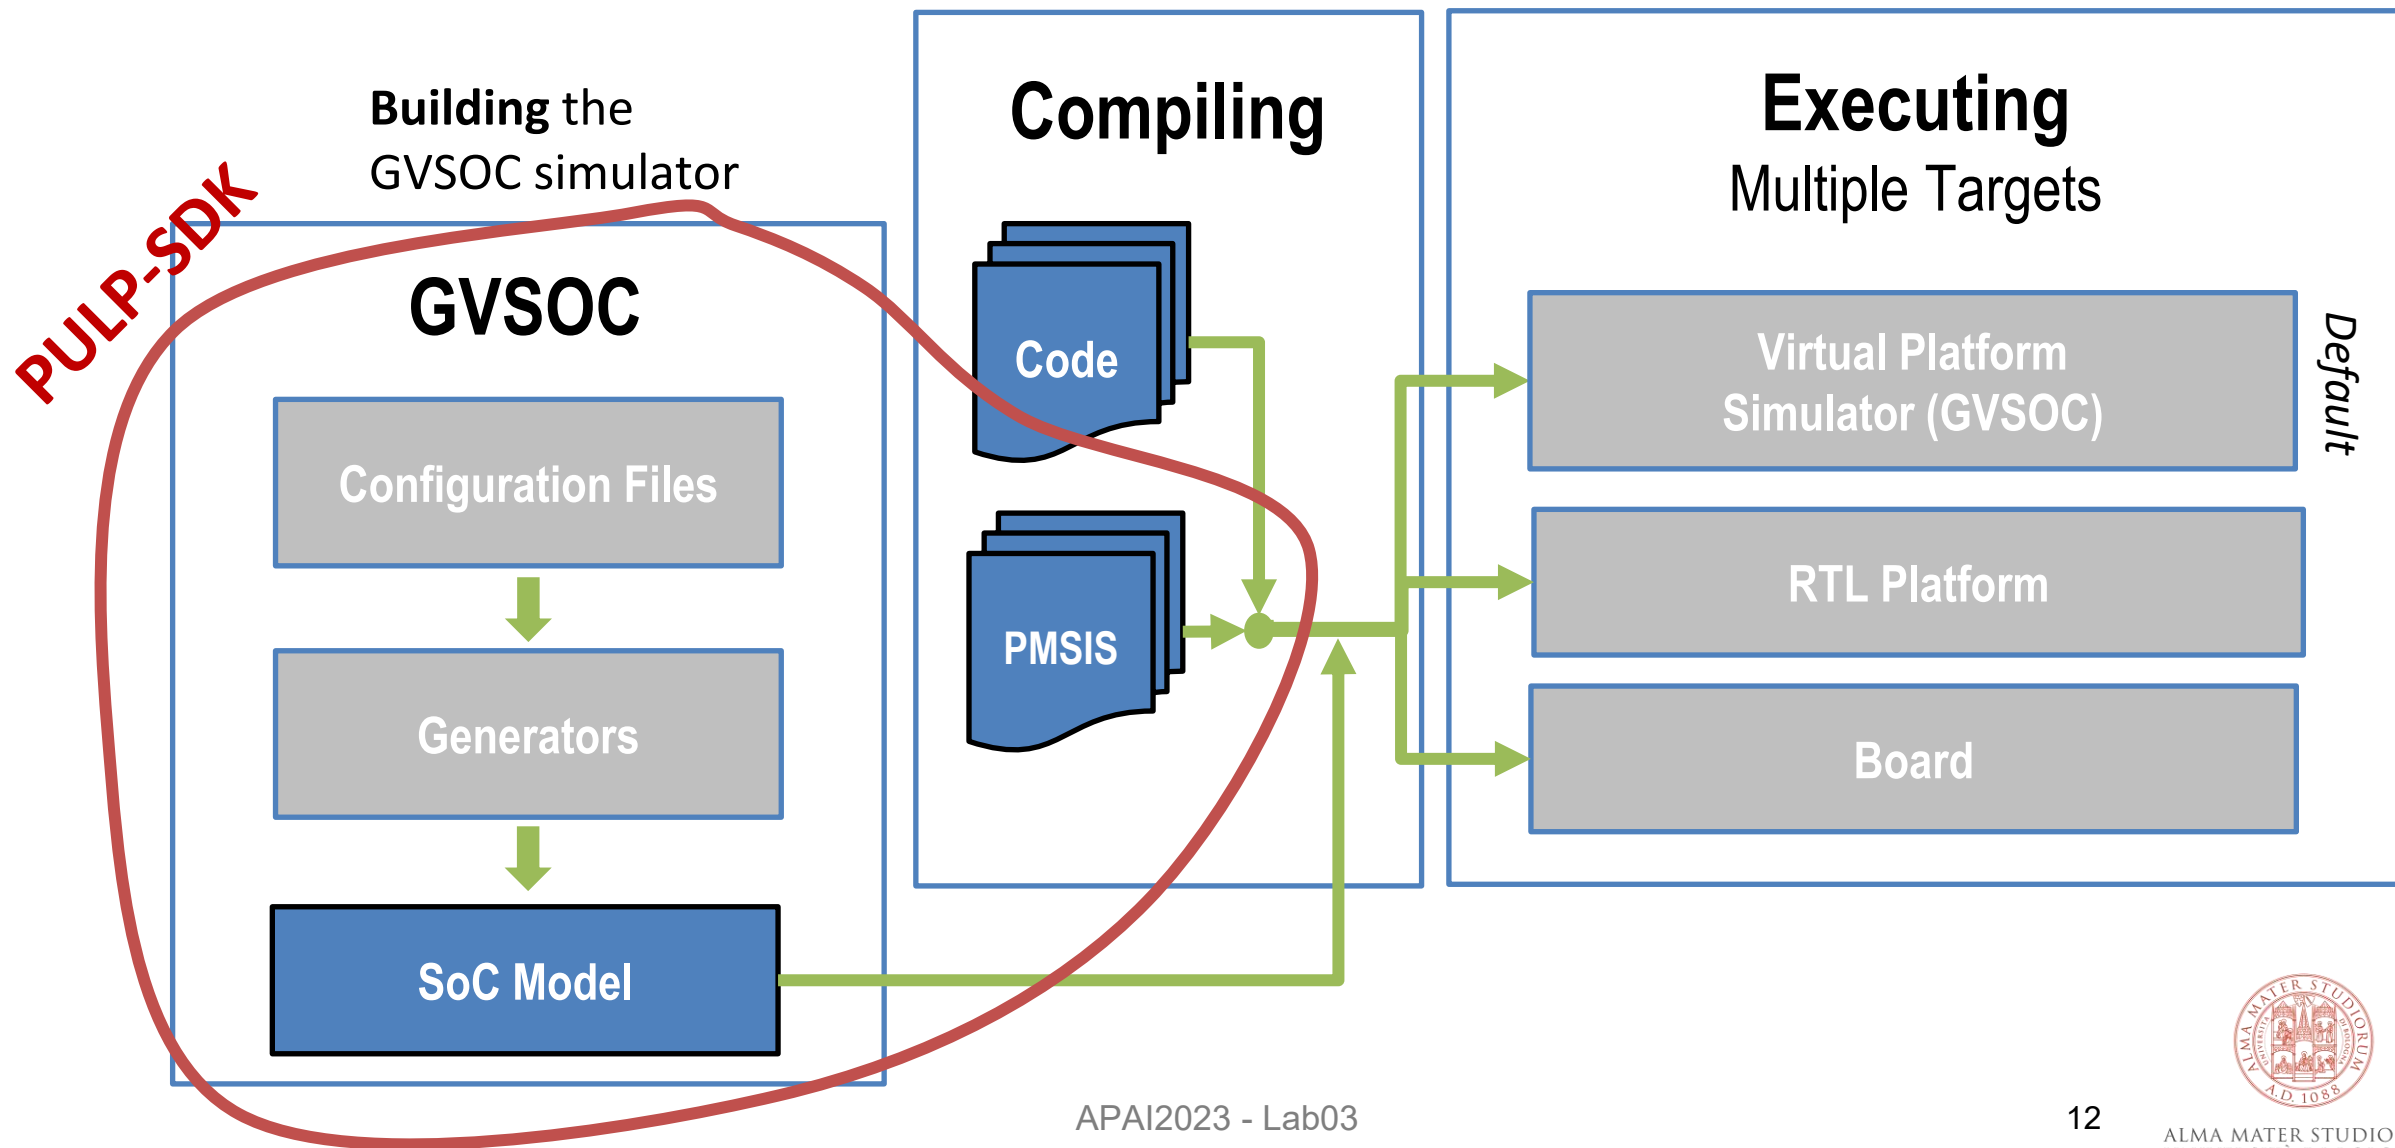
**Application code:** *main* function, application libraries, e.g. CNN inference function code...

**Runtime SW:** Peripheral *Drivers*, RTOS, Board support packages (BSP)

- *PMSIS layer*



# PULP Software Environment Workflow



# PULP SDK (software development kit)

The PULP-SDK (<https://github.com/pulp-platform/pulp-sdk>) includes a **PULP platform simulator (GVSOC)** and the SW libraries.

Check the `/pulp/pulp-sdk` folder:

- `rtos/` runtime code and software stack
- `tools/` configuration files, python generators, pulp runner and the gvsoc components
- `tests/` sample code to test the platform's features
- `applications/` full-application codes

```
$ cd /pulp/pulp-sdk
```

```
pulp-user@pulp-box /pulp/pulp-sdk $ ll
total 60
drwxrwxr-x  9 pulp-user pulp-user 4096 Feb  3  2021 ./
drwxrwxr-x  6 pulp-user pulp-user 4096 Feb  3  2021 ../
drwxrwxr-x  6 pulp-user pulp-user 4096 Feb  3  2021 build/
drwxrwxr-x  2 pulp-user pulp-user 4096 Feb  3  2021 configs/
drwxrwxr-x  8 pulp-user pulp-user 4096 Feb  3  2021 .git/
-rw-rw-r--  1 pulp-user pulp-user   80 Feb  3  2021 .gitignore
drwxrwxr-x  4 pulp-user pulp-user 4096 Feb  3  2021 install/
-rw-rw-r--  1 pulp-user pulp-user 11357 Feb  3  2021 LICENSE
-rw-rw-r--  1 pulp-user pulp-user   877 Feb  3  2021 Makefile
-rw-rw-r--  1 pulp-user pulp-user  2455 Feb  3  2021 README.md
drwxrwxr-x  4 pulp-user pulp-user 4096 Feb  3  2021 rtos/
drwxrwxr-x  2 pulp-user pulp-user 4096 Feb  3  2021 rules/
drwxrwxr-x 10 pulp-user pulp-user 4096 Feb  3  2021 tools/
pulp-user@pulp-box /pulp/pulp-sdk $
```



You could get the latest update from the open-source community and rebuild the GVSOC simulator

```
$ git pull origin main
$ source configs/pulp-open.sh
$ make build
```

**Do NOT do it**

# GVSoc – Features

## Virtual platform features:

- *C++* for fast native simulation
- *Python* for instantiation + configuration
- Complete set of traces to see what happen

## Timing model:

- Fully-event based, instances can generate events at specific time
- Includes timing models for interconnects, DMACs, memories...
- Performance counters for information from the execution

## Simulation performance:

- Around 1MIPS simulation speed
- Functionally aligned and calibrated with HW
- Timing accuracy is within 10-20% of target HW



ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA

# TASK1

# TASK 1: vector sum example

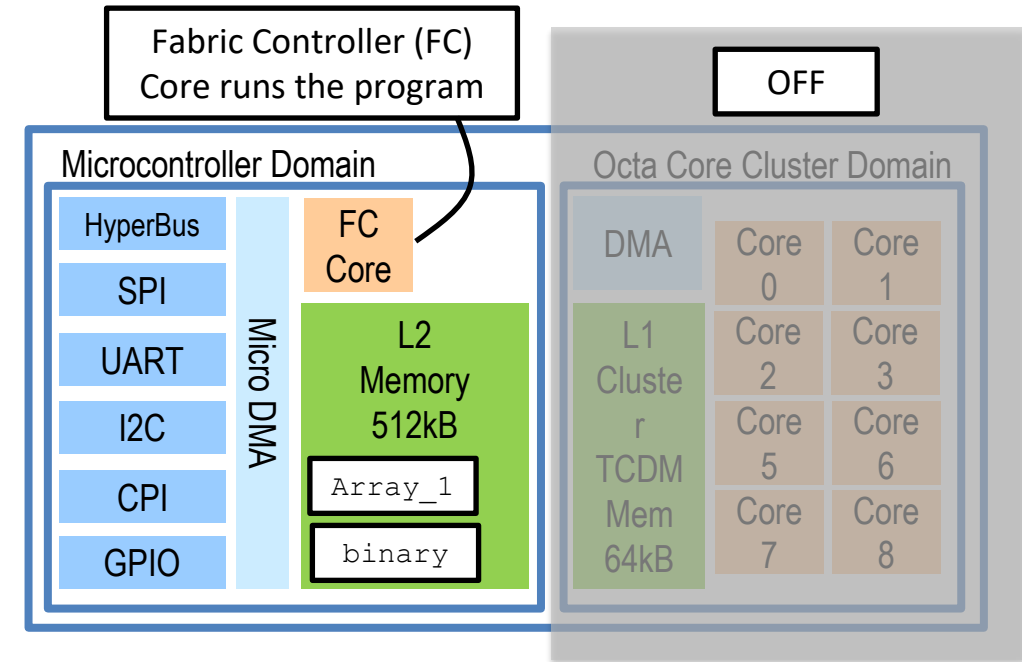
## Preparation

### 1. Clone repository for the lab, if you didn't already do it

```
git clone https://github.com/EEESlab/APAI25-LAB01-PULP-Embedded-Programming
```

### 2. Go into folder for task1

```
cd APAI25-LAB01-PULP_Embedded_Programming/  
cd vector_sum/  
make clean all run
```



## Tasks

### Read your assignment!

The `vector_sum()` function returns the element-wise add of the values in `array_1[N]`.

Note: the `main` include the function

$$\text{testbench} : S = \sum_{i=1}^N i = \frac{N \cdot (N-1)}{2}$$

Good practice for test!

- I. Array initialization ( $a[i] = i$ )
- II. Function Call
- III. Check Result

Task  
1.2

Take now a look to the code.

**What happen if you change:**

(line 13) `#define N 350`

➤ **How would you solve it?**

APAI2023 -



ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA



## SOLUTION: Casting Variables

(line 13) `#define N 350`

*Output:* Result is **not** correct. Got 37105 instead of 61425

If printing the `array_1` values after initialization:

```
1 2 3 4 5 6 7 8 9 10 11 12 13 14 [...] 247 248 249 250 251 252 253 254 255 0 1 2 3 4 5 6 7 8 [...] 88 89 90 91 92 93 94
```

- `unsigned char` datatype can represent numbers from 0 to 255!! If  $>255$ , assigned values get truncated ( $\text{value} \% 256$ )!
- **Solution:** cast `array_1` to `int` or `short int` (and the function's arguments!)

## (optional) Passing parameters via `Makefile`

(OPTIONAL) After applying the fix

We can pass parameters via Makefile:  
comment line 13, change the *Makefile* and  
launch:

```
$ make clean all run N=350
```

APP = test

N?=50

APP\_SRCS += test.c

APP\_CFLAGS += -O3 -g

APP\_CFLAGS += -DN=\$(N)

APP\_LDFLAGS +=

include \$(RULES\_DIR)/pmsis\_rules.mk

*Makefile*






ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA

# Intro: TASK2


## TASK2: Matrix-Vector Product



```
$ cd matrix_vector/  
$ make clean all run
```

## Example II: Matrix-Vector Product


A multi-dimensional array, including a matrix, is efficiently represented as a memory-contiguous array (**NOT**  $\text{Mat}[M][N]$ ).



In general:  $\text{Mat}[i][j] \rightarrow M[i*M+j]$

```
$ cd matrix_vector/  
$ make clean all run
```

## Example II: Matrix-Vector Product



```
// generic matrix-vector multiplication
int gemv(int N, int M, float * mat, float *vec, float * output_vec){

    for (int i=0; i<N; i++){
        for (int j=0; j<M; j++){
            vec_o[i] += mat_i[i*size_M+j] * vec_i[j];
        }
    }

}
```

# Assembly Code

A disassembler is a computer program that translates machine language into assembly language → the inverse operation to that of an assembler.

**Disassembly**, the output of a disassembler, is often formatted for human-readability rather than suitability for input to an assembler, making it principally a reverse-engineering tool. *[Wikipedia]*

To obtain the assembly code use the command:

```
$ make dis > dis.txt
```

Check the RI5CY User manual:

<https://riscv.org/wp-content/uploads/2017/05/riscv-spec-v2.2.pdf>

[https://www.pulp-platform.org/docs/ri5cy\\_user\\_manual.pdf](https://www.pulp-platform.org/docs/ri5cy_user_manual.pdf)

# Assembly Code

```

1c008706 <gemv>:
1c008706: 02a05d63      blez a0,1c008740 <gemv+0x3a>
1c00870a: 00259f13      slli t5,a1,0x2
1c00870e: 8fb6         mv t6,a3
1c008710: 01e68e33      add t3,a3,t5
1c008714: 4e81         li t4,0
1c008716: a005         j 1c008736 <gemv+0x30>
1c008718: 0048230b      p.lw t1,4(a6!)
1c00871c: 0046a88b      p.lw a7,4(a3!)
1c008720: 431c         lw a5,0(a4)
1c008722: 431307b3      p.mac a5,t1,a7
1c008726: c31c         sw a5,0(a4)
1c008728: ffc698e3      bne a3,t3,1c008718 <gemv+0x12>
1c00872c: 0e85         addi t4,t4,1
1c00872e: 967a         add a2,a2,t5
1c008730: 0711         addi a4,a4,4
1c008732: 01d50763      beq a0,t4,1c008740 <gemv+0x3a>
1c008736: 86fe         mv a3,t6
1c008738: 8832         mv a6,a2
1c00873a: fcb04fe3      bgtz a1,1c008718 <gemv+0x12>
1c00873c: b7fd         j 1c00872c <gemv+0x26>

```

1c008738: 8832	mv a6,a2
PC	Instruction

Program  
counter

```

// generic matrix-vector multiplication
void gemv(int size_N, int size_M,
          float * mat_I, float *vec_i, float * vec_o)
{
    for (int i=0; i<size_N; i++){
        for (int j=0; j<size_M; j++){
            // multiply accumulate operation
            vec_o[i] += mat_I[i*size_M+j] * vec_i[j];
        }
    }
}

```

Inner  
loop

In the **Makefile**:

1. Change from -O1 to -O3. What has changed in the assembly?
2. Remove -mnohwloop. What has changed in the assembly?

NB. After changing the compiler flags, compile the code again and generate and disassembly using `make dis`

# Assembly Code: O3 optimization

APP = matrix-vector

APP\_SRCS += test.c

APP\_CFLAGS += **-O3** -g -mnohwloops

APP\_LDFLAGS +=

include \$(RULES\_DIR)/pmsis\_rules.mk

```
1c008706 <gemv>:
1c008706: 02a05c63      blez a0,1c00873e <gemv+0x38>
1c00870a: 02b05a63      blez a1,1c00873e <gemv+0x38>
1c00870e: 00259e93      slli t4,a1,0x2
1c008712: 050a         slli a0,a0,0x2
1c008714: 00a70e33      add t3,a4,a0
1c008718: 01d68333      add t1,a3,t4
1c00871c: 0047258b      p.lw a1,4(a4!) # 10004 <__l1_heap_size+0x20>
1c008720: 87b6         mv a5,a3
1c008722: 8532         mv a0,a2
1c008724: 0045288b      p.lw a7,4(a0!)
1c008728: 0047a80b      p.lw a6,4(a5!)
1c00872c: 430885b3      p.mac a1,a7,a6
1c008730: feb72e23      sw a1,-4(a4)
1c008734: fef318e3      bne t1,a5,1c008724 <gemv+0x1e>
1c008738: 9676         add a2,a2,t4
1c00873a: feee11e3      bne t3,a4,1c00871c <gemv+0x16>
1c00873e: 8082         ret
```

The -O3 optimized code get rid of a `lw` instruction because the accumulator is kept in the register file!

```
// generic matrix-vector multiplication
void gemv(int size_N, int size_M,
float * mat_I, float *vec_i, float * vec_o)
{
    for (int i=0; i<size_N; i++){
        for (int j=0; j<size_M; j++){
            // multitply accumulate operation
            vec_o[i] += mat_i[i*size_M+j] * vec_i[j];
        }
    }
}
```

Inner  
loop



# Assembly Code: HW loops

APP = matrix-vector

APP\_SRCS += test.c

APP\_CFLAGS += -O3 -g -mnohwloops

APP\_LDFLAGS +=

include \$(RULES\_DIR)/pmsis\_rules.mk

```
1c008706 <gemv>:
1c008706: 04a05463    blez      a0,1c00874e <ge
1c00870a: 04b05263    blez      a1,1c00874e <ge
1c00870e: 00251e13    slli      t3,a0,0x2
1c008712: 1e71        addi      t3,t3,-4
1c008714: 00259e93    slli      t4,a1,0x2
1c008718: 002e5e13    srli      t3,t3,0x2
1c00871c: 01d68f33    add       t5,a3,t4
1c008720: 0e05        addi      t3,t3,1
1c008722: 015e407b    lp.setup   x0,t3,1c00874c <gemv+0x46>
1c008726: 40df07b3    sub       a5,t5,a3
1c00872a: 17f1        addi      a5,a5,-4
1c00872c: 0047258b    p.lw      a1,4(a4!) # 10004 <__l1_heap_size+0x20>
1c008730: 8389        srli      a5,a5,0x2
1c008732: 8836        mv        a6,a3
1c008734: 8532        mv        a0,a2
1c008736: 0785        addi      a5,a5,1
1c008738: 0087c0fb    lp.setup   x1,a5,1c008748 <gemv+0x42>
1c00873c: 0045230b    p.lw      t1,4(a0!)
1c008740: 0048288b    p.lw      a7,4(a6!)
1c008744: 431305b3    p.mac     a1,t1,a7
1c008748: feb72e23    sw        a1,-4(a4)
1c00874c: 9676        add       a2,a2,t4
1c00874e: 8082        ret
```

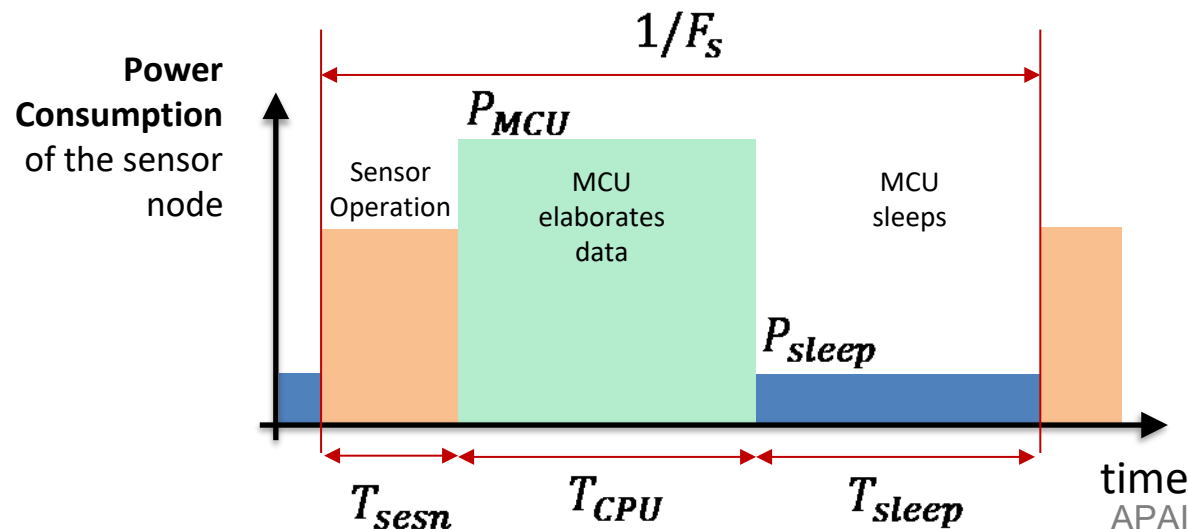

Hardware Loops Instructions have been placed!

Inner  
loop

```
// generic matrix-vector multiplication
void gemv(int size_N, int size_M,
float * mat_I, float *vec_i, float * vec_o)
{
    for (int i=0; i<size_N; i++){
        for (int j=0; j<size_M; j++){
            // multiply accumulate operation
            vec_o[i] += mat_i[i*size_M+j] * vec_i[j];
        }
    }
}
```



# Assessing the MCU performance



## Goal

- Extending the battery lifetime
  - Minimize the sensor node energy consumption
- Real time processing of sensor data

$$\min E_s + E_{MCU}$$

$$\text{s.t. } T_{CPU} < 1/F_s$$

Assuming:

- a fixed sample rate  $F_s$  a negligible sensor energy cost  $E_s \ll E_{MCU}$  (e.g.,  $T_{sesn} \ll 1/F_s$ )
- a constant power envelope of the MCU for active ( $P_{MCU}$ ) and sleep ( $P_{sleep} \ll P_{MCU}$ ) modes

$$\min T_{CPU} P_{MCU} + T_{sleep} P_{sleep}$$


$$T_{CPU} + T_{sleep} = 1/F_s$$

**Aim:** minimize the processing time  $T_{CPU}$

# What can we optimize? -> the execution time

The processing time  $T_{CPU}$  composes of:

- the **execution time**, where the CPU
- the **idle time**, where the CPU waits for events or IRQ (the core may be clock gated to save power)



Given a CPU clock frequency, the performance of a task (e.g. a C function) can be measured by accounting:

- the number of **elapsed clock cycles** ( $N_{clk}$ )
- the **number of instructions** ( $N_{instr}$ ) executed within the task depends on the code optimization!
- **CPI (Clock Cycles Per Instruction)**: different instructions may take a different number of **clock cycles** (depending on the CPU microarchitecture)

$$T_{CPU} = N_{clk} \times T_{clk} = CPI_{avg} \times N_{instr} \times T_{clk}$$

$$CPI_{avg} = \frac{N_{clk}}{N_{instr}}$$

$N_{clk}$  and  $N_{instr}$  greatly affects CPU performance!



# Performance Counters: Measuring the CPU Time

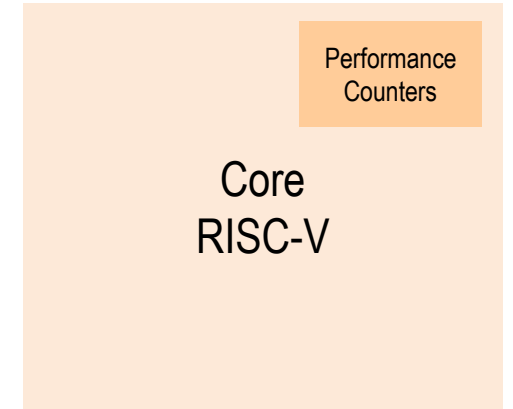
Each RI5CY cores of the PULP platform provide a **performance counter**. These 32-bit counters can be configured to count the:

- Total number of cycles (also includes the cycles where the core is sleeping)
- Number of cycles the core was active (not sleeping)
- Number of instructions executed
- Number of load data hazards
- Number of jump register data hazards
- Number of cycles waiting for instruction fetches, i.e. number of instructions wasted due to non-ideal caching
- Number of data memory loads executed. Misaligned accesses are counted twice
- Number of data memory stores executed. Misaligned accesses are counted twice
- Number of unconditional jumps (j, jal, jr, jalr)
- Number of both taken and not taken branches
- Number of taken branches
- Number of compressed instructions executed

The performance counters of the cluster cores can also account:

- Number of memory loads to EXT executed. Misaligned accesses are counted twice. Every non-L1 access is considered external
- Number of memory stores to EXT executed. Misaligned accesses are counted twice. Every non-L1 access is considered external
- Number of cycles used for memory loads to EXT. Every non-L1 access is considered external
- Number of cycles used for memory stores to EXT. Every non-L1 access is considered external
- Number of cycles wasted due to L1/log-interconnect contention

**Important:** while on the PULP Virtual Platform multiple perf counters can be enabled concurrently, only one counter is available on a real HW device!



# Using the performance Counters on PULP

Using the PMSIS library  
APIs:

Configure which performance  
counters to enable!

```
// enable the perf counters of interest
pi_perf_conf(    1 << PI_PERF_CYCLES |
                1 << PI_PERF_INSTR );

// reset the performance counters
pi_perf_reset();
// start the performance counters
pi_perf_start();

// task to profile
foo();

// stop the performance counters
pi_perf_stop();

// collect and print statistics
uint32_t instr_cnt = pi_perf_read(PI_PERF_INSTR);
uint32_t cycles_cnt = pi_perf_read(PI_PERF_CYCLES);
```

```
typedef enum {
    PI_PERF_CYCLES      = 17, /*!< Total number of cycles (also includes the
                                cycles where the core is sleeping). Be careful that this event is using a
                                timer shared within the cluster, so resetting, starting or stopping it on
                                one core will impact other cores of the same cluster. */
    PI_PERF_ACTIVE_CYCLES = 0, /*!< Counts the number of cycles the core was
                                active (not sleeping). */
    PI_PERF_INSTR       = 1, /*!< Counts the number of instructions executed.
                                */
    PI_PERF_LD_STALL    = 2, /*!< Number of load data hazards. */
    PI_PERF_JR_STALL    = 3, /*!< Number of jump register data hazards. */
    PI_PERF_IMISS       = 4, /*!< Cycles waiting for instruction fetches, i.e.
                                number of instructions wasted due to non-ideal caching. */
    PI_PERF_LD          = 5, /*!< Number of data memory loads executed.
                                Misaligned accesses are counted twice. */
    PI_PERF_ST          = 6, /*!< Number of data memory stores executed.
                                Misaligned accesses are counted twice. */
    PI_PERF_JUMP        = 7, /*!< Number of unconditional jumps (j, jal, jr,
                                jalr). */
    PI_PERF_BRANCH      = 8, /*!< Number of branches. Counts both taken and
                                not taken branches. */
    PI_PERF_BTAKEN      = 9, /*!< Number of taken branches. */
    PI_PERF_RVC         = 10, /*!< Number of compressed instructions
                                executed. */
    PI_PERF_LD_EXT      = 12, /*!< Number of memory loads to EXT executed.
                                Misaligned accesses are counted twice. Every non-TCDM access is considered
                                external (cluster only). */
    PI_PERF_ST_EXT      = 13, /*!< Number of memory stores to EXT executed.
                                Misaligned accesses are counted twice. Every non-TCDM access is considered
                                external (cluster only). */
    PI_PERF_LD_EXT_CYC  = 14, /*!< Cycles used for memory loads to EXT.
                                Every non-TCDM access is considered external (cluster only). */
    PI_PERF_ST_EXT_CYC  = 15, /*!< Cycles used for memory stores to EXT.
                                Every non-TCDM access is considered external (cluster only). */
    PI_PERF_TCDM_CONT    = 16, /*!< Cycles wasted due to TCDM/log-interconnect
                                contention (cluster only). */
} pi_perf_event_e;
```

/rtos/pmsis/pmsis\_api/include/pmsis/chips/default.h



ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA

# TASK2

# Profile the *gemv*

Make a new project by copying the *matrix-vector gemv* example into a new folder

```
$ cd matrix_vector/  
make clean all run
```

For any of the previous compiler optimization, **measure the performance** using the PMSIS performance counters and report:

- number of clock cycles.
- number of instructions.
- Number of Multiply-Add operations
- the CPI.
- compute the number of clock cycles and instructions per *elementary operation*.
  - Define 1 elementary operation == 1 Multiply-and-Accumulate

$$a \leftarrow a + (b \times c)$$

#MAC = N x M = 50 x 50 = 2500

	-O1	-O3	-O3 HWLoops
Clock Cycles			
Instr.			
MAC			
CPI			
Intr/Cycles			
Instr / MAC			

# SOLUTION: Profile the *gemv*

Make a new project by copying the *matrix-vector gemv* example into a new folder

```
$ cd matrix_vector/  
make clean all run
```

## Solution:

```
pi_perf_conf(1<<PI_PERF_CYCLES | 1<<PI_PERF_INSTR);  
pi_perf_reset();  
  
pi_perf_start();  
// call the matrix-vector fucntion  
gemv(N, M, matrix, vector, output_vec);  
pi_perf_stop();  
  
uint32_t instr_cnt = pi_perf_read(PI_PERF_INSTR);  
uint32_t cycles_cnt = pi_perf_read(PI_PERF_CYCLES);  
printf("Num.Istr: %d Num.Cycles: %d \n", instr_cnt,  
cycles_cnt);
```

#MAC = N x M = 50 x 50 = 2500

	-01	-03	-03 HWLoops
Clock Cycles			
Instr.			
MAC			
CPI			
Intr/Cycles			
Instr / MAC			





ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA